



Database Systems

Lecture 6: Advanced SQL

Dr. Momtazi
momtazi@aut.ac.ir

based on the slides of the course book



Outline

- **Accessing SQL From a Programming Language**
- Functions
- Triggers
- Advanced Aggregation Features
- OLAP



Accessing SQL From a Programming Language

- API (application-program interface) for a program to interact with a database server

- Application makes calls to
 - Connect with the database server
 - Send SQL commands to the database server
 - Fetch tuples of result one-by-one into program variables



Accessing SQL From a Programming Language

- Possible approaches:
 - Dynamic SQL
 - ▶ JDBC (Java Database Connectivity) works with Java
 - ▶ ODBC (Open Database Connectivity) works with C, C++, C#, and Visual Basic.
 - Other API's such as ADO.NET sit on top of ODBC
 - Embedded SQL



JDBC

- **JDBC** is a Java API for communicating with database systems supporting SQL.
- JDBC supports a variety of features for querying and updating data, and for retrieving query results.
- JDBC also supports metadata retrieval, such as querying about relations present in the database and the names and types of relation attributes.



JDBC

- Model for communicating with the database:
 - Open a connection
 - Create a “statement” object
 - Execute queries using the Statement object to send queries and fetch results
 - Exception mechanism to handle errors



JDBC Example

```
public static void JDBCexample(String userid, String passwd)
{
    try
    {
        Class.forName ("oracle.jdbc.driver.OracleDriver");
        Connection conn = DriverManager.getConnection(
            "jdbc:oracle:thin:@db.yale.edu:1521:univdb",
            userid, passwd);
        Statement stmt = conn.createStatement();
        try {
            stmt.executeUpdate(
                "insert into instructor values('77987', 'Kim', 'Physics', 98000)");
        } catch (SQLException sqle)
        {
            System.out.println("Could not insert tuple. " + sqle);
        }
        ResultSet rset = stmt.executeQuery(
            "select dept_name, avg (salary) "+
            " from instructor "+
            " group by dept_name");
        while (rset.next()) {
            System.out.println(rset.getString("dept_name") + " " +
                               rset.getFloat(2));
        }
        stmt.close();
        conn.close();
    }
    catch (Exception sqle)
    {
        System.out.println("Exception : " + sqle);
    }
}
```

Apago PDF Enhancer



JDBC Example

```
public static void JDBCexample(String userid, String passwd)
{
    try
    {
        Class.forName ("oracle.jdbc.driver.OracleDriver");
        Connection conn = DriverManager.getConnection(
            "jdbc:oracle:thin:@db.yale.edu:1521:univdb",
            userid, passwd);
        Statement stmt = conn.createStatement();
        try {
            stmt.executeUpdate(
                "insert into instructor values('77987', 'Kim', 'Physics', 98000)");
        } catch (SQLException sqle)
        {
            System.out.println("Could not insert tuple. " + sqle);
        }
    }
}
```




JDBC Example

```
ResultSet rset = stmt.executeQuery(
    "select dept_name, avg (salary) "+
    " from instructor "+
    " group by dept_name");
while (rset.next()) {
    System.out.println(rset.getString("dept_name") + " " +
        rset.getFloat(2));
}
stmt.close();
conn.close();
}
catch (Exception sqle)
{
    System.out.println("Exception : " + sqle);
}
}
```

Apago PDF Enhancer



Database Connection

- Each database product that supports JDBC provides a JDBC driver that must be dynamically loaded in order to access the database from Java.
 - This is done by invoking *Class.forName* with one argument specifying a concrete class implementing the `java.sql.Driver` interface
- Connecting to the Database: A connection is opened using the `getConnection` method of the `DriverManager` class (within `java.sql`) using 3 parameters:
 - a string that specifies the URL, or machine name, where the server runs
 - a database user identifier, which is a string
 - a password, which is also a string.
 - ▶ Note: the need to specify a password within the JDBC code presents a security risk if an unauthorized person accesses your Java code.



Shipping SQL Statements

- Methods for executing a statement:
 - *executeQuery*
 - ▶ When the SQL statement is a query
 - ▶ It returns a result set
 - *executeUpdate*
 - ▶ When the SQL statement is nonquery (DDL or DML)
 - Update
 - Insert
 - Delete
 - Create table
 - ...
 - ▶ It returns an integer giving the number of tuples inserted, updated, or deleted.
 - ▶ For DDL statements, the return value is zero.



Retrieving the Results of a Query

- Retrieving the set of tuples in the result into a `ResultSet` object
- Fetching the results one tuple at a time
- Using the *next* method on the result set to test whether there remains at least one unfetched tuple in the result set and if so, fetches it.
- Attributes from the fetched tuple are retrieved using various methods whose names begin with *get*
 - *getString*: can retrieve any of the basic SQL data types
 - *getFloat*
- Possible argument to the *get* methods
 - The attribute name specified as a string
 - An integer indicating the position of the desired attribute within the tuple



Database Connection

- The statement and connection are both closed at the end of the Java program.
- It is important to close the connection because there is a limit imposed on the number of connections to the database
- Unclosed connections may cause that limit to be exceeded.
- If this happens, the application cannot open any more connections to the database.



Prepared Statements

- Creating a prepared statement in which some values are replaced by “?”
- Specifying that actual values will be provided later
- Compiling the query by the database system when it is prepared
- Reusing the previously compiled form of the query and apply the new values whenever the query is executed
 - (with new values to replace the “?”s),



Prepared Statements

```
PreparedStatement pStmt = conn.prepareStatement(
    "insert into instructor values(?,?,?,?)");
pStmt.setString(1, "88877");
pStmt.setString(2, "Perry");
pStmt.setString(3, "Finance");
pStmt.setInt(4, 125000);
pStmt.executeUpdate();
pStmt.setString(1, "88878");
pStmt.executeUpdate();
```

Figure 5.2 Prepared statements in JDBC code.

- insert into instructor values ("88877", "Perry", "Finance", 125000)
- insert into instructor values ("88878", "Perry", "Finance", 125000)



Prepared Statements

- Prepared statements allow for more efficient execution
 - Where the same query can be compiled once and then run multiple times with different parameter values
 - Whenever a user-entered value is used, even if the query is to be run only once



- ```
ResultSetMetaData rsmd = rs.getMetaData();
for(int i = 1; i <= rsmd.getColumnCount(); i++) {
 System.out.println(rsmd.getColumnName(i));
 System.out.println(rsmd.getColumnTypeName(i));
}
```



# ODBC

- Open DataBase Connectivity (ODBC) standard
  - standard for application program to communicate with a database server.
  - application program interface (API) to
    - ▶ open a connection with a database,
    - ▶ send queries and updates,
    - ▶ get back results.



# Embedded SQL

- The SQL standard defines embeddings of SQL in a variety of programming languages such as C, C++, Java, Fortran, and PL/1,
- A language to which SQL queries are embedded is referred to as a **host language**
- The SQL structures permitted in the host language comprise *embedded SQL*
- An embedded SQL program must be processed by a special preprocessor prior to compilation.
- The preprocessor replaces embedded SQL requests with host-language declarations and procedure calls that allow runtime execution of the database accesses.
- Then, the resulting program is compiled by the host-language compiler.



# Embedded SQL

- **EXEC SQL** statement is used to identify embedded SQL request to the preprocessor

EXEC SQL <embedded SQL statement >;

- Note: this varies by language:
  - In some languages, like COBOL, the semicolon is replaced with END-EXEC
  - In Java embedding uses  
# SQL { .... };



# Database Connection

- Before executing any SQL statements, the program must first connect to the database. This is done using:

EXEC-SQL **connect to** *server* **user** *user-name* **using** *password*;

- Here, *server* identifies the server to which a connection is to be established.



# Variables

- Variables of the host language can be used within embedded SQL statements.
- They are preceded by a colon (:) to distinguish from SQL variables (e.g., *:credit\_amount* )
- Variables used as above must be declared within DECLARE section, as illustrated below. The syntax for declaring the variables, however, follows the usual host language syntax.

```
EXEC-SQL BEGIN DECLARE SECTION;
```

```
 int credit_amount ;
```

```
EXEC-SQL END DECLARE SECTION;
```



# SQL Query

- To write an embedded SQL query, we use the following statement:

**declare c cursor for <SQL query>**

- The variable *c* is used to identify the query

- Example:

**EXEC SQL**

**declare c cursor for**

**select *ID, name***

**from *student***

**where *tot\_cred* > *:credit\_amount***

**END\_EXEC**



# SQL Query

- The open statement is then used to evaluate the query
- The **open** statement for our example is as follows:

**EXEC SQL open c ;**

- This statement causes the database system to execute the query and to save the results within a temporary relation.
- The query uses the value of the host-language variable *credit-amount* at the time the **open** statement is executed.





# SQL Query

## ■ The fetch statement

- Placing the values of one tuple in the query result into host language variables
- Requiring one host-language variable for each attribute of the result relation;

e.g., we need one variable to hold the ID value (si) and another to hold the name value (sn) which have been declared within a DECLARE section

**EXEC SQL**

**fetch c into :si, :sn**

**END\_EXEC**

Repeated calls to fetch get successive tuples in the query result



# SQL Query

- The **close** statement causes the database system to delete the temporary relation that holds the result of the query.

**EXEC SQL close c ;**



# Outline

- Accessing SQL From a Programming Language
- **Functions**
- Triggers
- Advanced Aggregation Features
- OLAP



# Functions and Procedures

- SQL:1999 supports functions and procedures
  - Functions/procedures can be written in SQL itself, or in an external programming language (e.g., C, Java).
  - Some database systems support **table-valued functions**, which can return a relation as a result.
- SQL:1999 also supports a rich set of imperative constructs, including
  - Loops, if-then-else, assignment
- Many databases have proprietary procedural extensions to SQL that differ from SQL:1999.



# SQL Functions

- Define a function that, given the name of a department, returns the count of the number of instructors in that department.

```
create function dept_count (dept_name varchar(20))
 returns integer
 begin
 declare d_count integer;
 select count (*) into d_count
 from instructor
 where instructor.dept_name = dept_name
 return d_count;
 end
```



# SQL Functions

- The function *dept\_count* can be used in a query that returns names and budget of all departments with more that 12 instructors.

```
select dept_name, budget
from department
where dept_count (dept_name) > 12
```



# SQL Functions

- Compound statement: **begin ... end**
  - May contain multiple SQL statements between **begin** and **end**.
- **returns** -- indicates the variable-type that is returned (e.g., integer)
- **return** -- specifies the values that are to be returned as result of invoking the function
- SQL function are in fact **parameterized views** that generalize the regular notion of views by allowing parameters.



# Table Functions

- SQL:2003 added functions that return a relation as a result
- Example: Return all instructors in a given department

**create function** *instructor\_of* (*dept\_name* **char**(20))

**returns table** (

*ID* **varchar**(5),  
*name* **varchar**(20),  
*dept\_name* **varchar**(20),  
*salary* **numeric**(8,2))

**return table**

(**select** *ID*, *name*, *dept\_name*, *salary*  
**from** *instructor*  
**where** *instructor.dept\_name* = *instructor\_of.dept\_name*)

- Usage

**select** \*  
**from table** (*instructor\_of* ('Music'))





# Outline

- Accessing SQL From a Programming Language
- Functions
- **Triggers**
- Advanced Aggregation Features
- OLAP



# Triggers

- A **trigger** is a statement that is executed automatically by the system as a side effect of a modification to the database.
- To design a trigger mechanism, we must:
  - Specify the conditions under which the trigger is to be executed.
  - Specify the actions to be taken when the trigger executes.
- Triggers introduced to SQL standard in SQL:1999, but supported even earlier using non-standard syntax by most databases.
  - Syntax illustrated here may not work exactly on your database system; check the system manuals



# Triggering Events and Actions in SQL

- Triggering event can be **insert**, **delete** or **update**
- Triggers on update can be restricted to specific attributes
  - For example, **after update of *takes on grade***
- Values of attributes before and after an update can be referenced
  - **referencing old row as** : for deletes and updates
  - **referencing new row as** : for inserts and updates
- Triggers can be activated before an event, which can serve as extra constraints. For example, convert blank grades to null.



# Trigger Example: to Maintain Referential Integrity

```
create trigger timeslot_check1 after insert on section
referencing new row as nrow
for each row
when (nrow.time_slot_id not in (
 select time_slot_id
 from time_slot)) /* time_slot_id not present in time_slot */
begin
 rollback
end;
```

```
create trigger timeslot_check2 after delete on timeslot
referencing old row as orow
for each row
when (orow.time_slot_id not in (
 select time_slot_id
 from time_slot) /* last tuple for time_slot_id deleted from time_slot */
and orow.time_slot_id in (
 select time_slot_id
 from section)) /* and time_slot_id still referenced from section */
begin
 rollback
end;
```



# Trigger Example: to Maintain `credits_earned` value

```
create trigger credits_earned after update of takes on (grade)
referencing new row as nrow
referencing old row as orow
for each row
when nrow.grade <> 'F' and nrow.grade is not null
 and (orow.grade = 'F' or orow.grade is null)
begin atomic
 update student
 set tot_cred = tot_cred +
 (select credits
 from course
 where course.course_id = nrow.course_id)
 where student.id = nrow.id;
end;
```



# Disabling Triggers

- Triggers can be disabled or enabled;
  - by default they are enabled when they are created.
- Triggers can be disabled by:  
**alter trigger *trigger\_name* disable**  
**disable trigger *trigger\_name***
- A trigger that has been disabled can be enabled again.
- A trigger can instead be dropped, which removes it permanently, by:  
**drop trigger *trigger\_name***



# When Not To Use Triggers

- Triggers were used earlier for tasks such as
  - Maintaining summary data (e.g., total salary of each department)
  - Replicating databases by recording changes to special relations (called **change** or **delta** relations) and having a separate process that applies the changes over to a replica
- There are better ways of doing these now:
  - Databases today provide built in materialized view facilities to maintain summary data
  - Databases provide built-in support for replication
- Encapsulation facilities can be used instead of triggers in many cases
  - Define methods to update fields
  - Carry out actions as part of the update methods instead of through a trigger



# When Not To Use Triggers

- Risk of unintended execution of triggers, for example, when
  - Loading data from a backup copy
  - Replicating updates at a remote site
  - Trigger execution can be disabled before such actions.
- Other risks with triggers:
  - Error leading to failure of critical transactions that set off the trigger
  - Cascading execution





# Outline

- Accessing SQL From a Programming Language
- Functions
- Triggers
- **Advanced Aggregation Features**
- OLAP



# Ranking

- Suppose we are given a relation  
*student\_grades*(*ID*, *GPA*)  
giving the grade-point average of each student
- Goal: finding the rank of each student.
- Ranking can be done using basic SQL aggregation, but resultant query is very inefficient

```
select ID, (1 + (select count(*)
 from student_grades B
 where B.GPA > A.GPA)) as s_rank
from student_grades A
order by s_rank;
```



# Ranking

- Ranking is done in conjunction with an order by specification.

```
select ID, rank() over (order by GPA desc) as s_rank
from student_grades
order by s_rank
```

- NOTE: the extra **order by** clause is needed to get them in sorted order



# Ranking

- Two possible approaches for ranking
  - Leaving gaps: (default)  
e.g. if 2 students have the same top GPA, both have rank 1, and the next rank is 3
  - Without gaps: (using **dense\_rank**)  
so next dense rank would be 2



# Ranking with Partitions

- Ranking can be done within partition of the data.
- “Find the rank of students within each department.”

```
select ID, dept_name,
 rank () over (partition by dept_name order by GPA desc)
 as dept_rank
from dept_grades
order by dept_name, dept_rank;
```

- Ranking is done *after* applying **group by** clause/aggregation



# Top $n$ Items

- Can be used to find top- $n$  results
  - More general than the **limit**  $n$  clause supported by many databases, since it allows top- $n$  within each partition



# Other Ranking Functions

- SQL:1999 permits the user to specify **nulls first** or **nulls last**

```
select ID,
 rank () over (order by GPA desc nulls last) as s_rank
from student_grades
```



# Outline

- Accessing SQL From a Programming Language
- Functions
- Triggers
- Advanced Aggregation Features
- **OLAP**





# Data Analysis and OLAP

## ■ Online Analytical Processing (OLAP)

- Interactive analysis of data, allowing data to be summarized and viewed in different ways in an online fashion (with negligible delay)
- Is used for **multidimensional data** (data that can be modeled as dimension attributes and measure attributes)
  - **Measure attributes**
    - ▶ measure some value
    - ▶ can be aggregated upon
    - ▶ e.g., the attribute *number* of the *sales* relation
  - **Dimension attributes**
    - ▶ define the dimensions on which measure attributes (or aggregates there of) are viewed
    - ▶ e.g., attributes *item\_name*, *color*, and *size* of the *sales* relation



## Example sales relation

| <i>item_name</i> | <i>color</i> | <i>clothes_size</i> | <i>quantity</i> |
|------------------|--------------|---------------------|-----------------|
| skirt            | dark         | small               | 2               |
| skirt            | dark         | medium              | 5               |
| skirt            | dark         | large               | 1               |
| skirt            | pastel       | small               | 11              |
| skirt            | pastel       | medium              | 9               |
| skirt            | pastel       | large               | 15              |
| skirt            | white        | small               | 2               |
| skirt            | white        | medium              | 5               |
| skirt            | white        | large               | 3               |
| dress            | dark         | small               | 2               |
| dress            | dark         | medium              | 6               |
| dress            | dark         | large               | 12              |
| dress            | pastel       | small               | 4               |
| dress            | pastel       | medium              | 3               |
| dress            | pastel       | large               | 3               |
| dress            | white        | small               | 2               |
| dress            | white        | medium              | 3               |
| dress            | white        | large               | 0               |
| shirt            | dark         | small               | 2               |
| shirt            | dark         | medium              | 6               |
| ...              | ...          | ...                 | ...             |
| ...              | ...          | ...                 | ...             |



# Cross Tabulation of sales by *item\_name* and color

*clothes\_size* **all**


|                  |       | <i>color</i> |        |       |       |
|------------------|-------|--------------|--------|-------|-------|
|                  |       | dark         | pastel | white | total |
| <i>item_name</i> | skirt | 8            | 35     | 10    | 53    |
|                  | dress | 20           | 10     | 5     | 35    |
|                  | shirt | 14           | 7      | 28    | 49    |
|                  | pants | 20           | 2      | 5     | 27    |
|                  | total | 62           | 54     | 48    | 164   |

- The table above is an example of a **cross-tabulation** (**cross-tab**), also referred to as a **pivot-table**.
  - Values for one of the dimension attributes form the row headers
  - Values for another dimension attribute form the column headers
  - Other dimension attributes are listed on top
  - Values in individual cells are (aggregates of) the values of the dimension attributes that specify the cell.



# Data Cube

- A **data cube** is a multidimensional generalization of a cross-tab
- Can have  $n$  dimensions; we show 3 below
- Cross-tabs can be used as views on a data cube





# Online Analytical Processing Operations

- **Pivoting:** changing the dimensions used in a cross-tab
  - Each cross-tab is a two-dimensional view on a multidimensional data cube.
  - With an OLAP system, a data analyst can look at different cross-tabs on the same data by interactively selecting the attributes in the cross-tab.
  - Examples:
    - ▶ selecting a cross-tab on item name and clothes size
    - ▶ selecting a cross-tab on color and clothes size



# Online Analytical Processing Operations

- **Slicing:** creating a cross-tab for fixed values only
  - OLAP systems allow an analyst to see a cross-tab on item name and color for a fixed value of clothes size,
  - Example:
    - ▶ large, instead of the sum across all sizes.
  - This operation is referred to as slicing, since it can be thought of as viewing a slice of the data cube.
  - The operation is sometimes called **dicing**.



# Cross Tabulation With Hierarchy

- The following table can be achieved using natural join with another table specifying item\_names and category

*clothes\_size:* **all**

| <i>category</i> | <i>item_name</i> | <i>color</i> |        |       |       |     |
|-----------------|------------------|--------------|--------|-------|-------|-----|
|                 |                  | dark         | pastel | white | total |     |
| womenswear      | skirt            | 8            | 8      | 10    | 53    | 88  |
|                 | dress            | 20           | 20     | 5     | 35    |     |
|                 | subtotal         | 28           | 28     | 15    |       |     |
| menswear        | pants            | 14           | 14     | 28    | 49    | 76  |
|                 | shirt            | 20           | 20     | 5     | 27    |     |
|                 | subtotal         | 34           | 34     | 33    |       |     |
| total           |                  | 62           | 62     | 48    |       | 164 |



# Cross Tabulation With Hierarchy

- Cross-tabs can be easily extended to deal with hierarchies
- Can drill down or roll up on a hierarchy
  - roll up: moving from finer granularity to coarser-granularity (by the mean of aggregation)
  - drill down: moving from coarser-granularity to finer granularity (must be generated from original data)





# Relational Representation of Cross-tabs

- Cross-tabs can be represented as relations
  - The value **all** is used to represent aggregates.
  - The SQL standard actually uses null values in place of **all** despite confusion with regular null values.

| <i>item_name</i> | <i>color</i> | <i>clothes_size</i> | <i>quantity</i> |
|------------------|--------------|---------------------|-----------------|
| skirt            | dark         | <b>all</b>          | 8               |
| skirt            | pastel       | <b>all</b>          | 35              |
| skirt            | white        | <b>all</b>          | 10              |
| skirt            | <b>all</b>   | <b>all</b>          | 53              |
| dress            | dark         | <b>all</b>          | 20              |
| dress            | pastel       | <b>all</b>          | 10              |
| dress            | white        | <b>all</b>          | 5               |
| dress            | <b>all</b>   | <b>all</b>          | 35              |
| shirt            | dark         | <b>all</b>          | 14              |
| shirt            | pastel       | <b>all</b>          | 7               |
| shirt            | White        | <b>all</b>          | 28              |
| shirt            | <b>all</b>   | <b>all</b>          | 49              |
| pant             | dark         | <b>all</b>          | 20              |
| pant             | pastel       | <b>all</b>          | 2               |
| pant             | white        | <b>all</b>          | 5               |
| pant             | <b>all</b>   | <b>all</b>          | 27              |
| <b>all</b>       | dark         | <b>all</b>          | 62              |
| <b>all</b>       | pastel       | <b>all</b>          | 54              |
| <b>all</b>       | white        | <b>all</b>          | 48              |
| <b>all</b>       | <b>all</b>   | <b>all</b>          | 164             |



# Extended Aggregation to Support OLAP

- The **cube** operation computes union of **group by**'s on every subset of the specified attributes
- Example relation for this section  
*sales(item\_name, color, clothes\_size, quantity)*
- E.g. consider the query

```
select item_name, color, size, sum(number)
from sales
group by cube(item_name, color, size)
```

This computes the union of eight different groupings of the *sales* relation:

```
{ (item_name, color, size),
 (item_name, color), (item_name, size), (color, size),
 (item_name), (color), (size), () }
```

where ( ) denotes an empty **group by** list.

- For each grouping, the result contains the null value for attributes not present in the grouping.



# Extended Aggregation

- The **rollup** construct generates union on every prefix of specified list of attributes
- E.g.,

```
select item_name, color, size, sum(number)
from sales
group by rollup(item_name, color, size)
```

Generates union of four groupings:

```
{ (item_name, color, size),
 (item_name, color),
 (item_name),
 () }
```



# Extended Aggregation

- Rollup can be used to generate aggregates at multiple levels of a hierarchy.
- E.g., suppose table *itemcategory*(*item\_name*, *category*) gives the category of each item. Then

```
select category, item_name, sum(number)
from sales, itemcategory
where sales.item_name = itemcategory.item_name
group by rollup(category, item_name)
```

would give a hierarchical summary by *item\_name* and by *category*.



# Extended Aggregation

- Multiple rollups or cubes can be used in a single group by clause
  - Each generates set of group by lists, cross product of sets gives overall set of group by lists

■ E.g.,

```
select item_name, color, size, sum(number)
from sales
group by rollup(item_name), rollup(color, size)
```

generates the groupings

$$\{item\_name, ()\} \times \{(color, size), (color), ()\}$$
$$= \{ (item\_name, color, size), (item\_name, color), (item\_name), (color, size), (color), ( ) \}$$



# OLAP Implementation

- The earliest OLAP systems used multidimensional arrays in memory to store data cubes, and are referred to as **multidimensional OLAP (MOLAP)** systems.
- OLAP implementations using only relational database features are called **relational OLAP (ROLAP)** systems
- Hybrid systems, which store some summaries in memory and store the base data and other summaries in a relational database, are called **hybrid OLAP (HOLAP)** systems.



# OLAP Implementation

- Early OLAP systems precomputed *all* possible aggregates in order to provide online response
  - Space and time requirements for doing so can be very high
    - ▶  $2^n$  combinations of **group by**
- Several optimizations available for computing multiple aggregates
  - It suffices to precompute some aggregates, and compute others on demand from one of the precomputed aggregates
    - ▶ Can compute aggregate on *(item\_name, color)* from an aggregate on *(item\_name, color, size)*
      - is cheaper than computing it from scratch



Questions?